



TOM NASH
THINK. ACT. REACT.

The Top 15 Investing Opportunities of the Decade

A Framework for Ownership in the Age of AI

By Tom Nash | December 18, 2025

Introduction: A Framework for Ownership

Hey it's Tom.

This document is built around one simple idea. The next 10 years will reward ownership of systems, platforms, and infrastructure far more than short lived narratives. Markets love drama. Wealth gets built by businesses that quietly become essential, then keep collecting cash while everyone argues about the latest headline.

Over the next decade, AI shifts from a feature you toggle on to a foundation you build everything on. When that happens, value concentrates in the layers underneath the excitement. Power generation that can run around the clock. Data centers that can scale. Chips that train and serve models. Cloud platforms that host the workloads. Security that keeps the whole machine from getting hacked into a smoking crater.

Software that turns messy real world operations into decisions and outcomes. These layers behave like dependencies, once they are embedded, removing them becomes painful, expensive, and sometimes impossible.

That is why this list looks the way it does. Many of these companies sell to enterprises and governments. They show up inside budgets, procurement cycles, and multi year contracts. They build products that become habits, workflows, and standards. Over time, that creates switching costs, pricing power, and durability. Durability matters because a 10 year holding period rewards businesses that keep compounding through multiple cycles, not businesses that win a popularity contest for 6 months.

A 10 year mindset changes the scoreboard. Quarterly noise fades into the background. Macro debates come and go. Management execution stays. Balance sheet strength stays. Product usefulness stays. The central question becomes simple. Does the business become harder to replace as the world becomes more complex and more automated. Every company on this list has a serious claim to that.

This is also a document about restraint. The goal here is not to chase miracles. The goal is to own compounding machines. Some names will look boring for long stretches. Some will disappoint for a quarter or two. Some will trade sideways for a year while the internet declares them “dead.” That experience is normal. Compounding rarely looks pretty in real time.

This is a framework for ownership. Build positions patiently. Add when fear is high and confidence is low. Keep your time horizon longer than the market’s attention span. Most investors lose because they overreact to noise and underweight the only thing that truly matters: time.

The next chapters break down each company one by one. What layer it owns, why that layer matters, how the moat gets stronger, and where the risks live. If you understand how these businesses fit together from energy to silicon to cloud to security to operating systems, you stop reacting to headlines and start thinking in systems. That shift alone can change your entire investing life.

The 15 Core Holdings

Palantir sits in a category most investors still misunderstand. It does not compete to build the smartest model or the flashiest demo. It exists to make sense of reality. That distinction matters because reality is messy, fragmented, political, regulated, and full of legacy systems that do not talk to each other.

At its core, Palantir turns raw data into decisions. Governments, militaries, hospitals, manufacturers, and large enterprises use it to connect data sources that were never designed to work together. Once those connections exist, Palantir sits in the middle of the workflow. Planning. Forecasting. Resource allocation. Risk management. This is where software becomes operational rather than decorative.

The reason Palantir scales so well in an AI driven world comes down to positioning. Models change. Chips change. Vendors rotate. Palantir stays above that fight. It plugs into whatever model a customer chooses and translates output into action. That neutrality gives it durability. If one model loses relevance and another takes over, Palantir keeps running. The value comes from orchestration, not prediction.

This is why the operating system analogy fits. Palantir does not replace existing systems. It sits on top of them. ERP systems. Databases. Sensors. Feeds. Human inputs. Once embedded, removing it becomes disruptive. Decisions slow down. Visibility disappears. Teams revert to spreadsheets and guesswork. That pain creates switching costs that do not show up on a balance sheet but absolutely show up in renewal rates.

Government contracts get the headlines, but the real story sits in commercial expansion. Manufacturing optimization. Supply chain resilience. Energy forecasting. Healthcare logistics. Financial crime detection. These problems grow as systems grow. Complexity works in Palantir's favor because complexity increases the cost of bad decisions.

The business model reflects that strength. High gross margins. Growing operating leverage. Customers expanding usage over time rather than churning out. This creates a flywheel where implementation leads to adoption, adoption leads to dependency, and dependency leads to durable cash flows.

There are risks. Sales cycles can be long. Public perception still lags reality. The stock will remain volatile because the story does not fit neatly into old valuation buckets. None of those issues break the long term thesis. They create noise, not structural weakness.

Over a ten year window, Palantir benefits from a simple truth. As organizations rely more on automated decision making, they need a layer that translates data into outcomes safely, audibly, and at scale. That layer does not need to win the AI race. It needs to work no matter who does. That is why Palantir belongs at the foundation of this list. It does not bet on a single future. It monetizes all of them.

Tesla \$TSLA

Autonomy, Robotaxis, and Robotics

Tesla is widely misunderstood because most people still look at it through the lens of a car company. That framing misses where the value is compounding. Vehicles are the distribution layer. The real assets sit in software, data, and autonomy.

Full self driving changes the equation completely. Every mile driven feeds a feedback loop that improves perception, planning, and decision making. That data advantage compounds over time. As the system improves, the marginal cost of deploying autonomy trends toward zero while the economic value of each deployment increases. Few businesses ever reach that dynamic.

The robotaxi opportunity flows naturally from this. Once autonomy reaches reliability at scale, the cost structure of transportation resets. No driver. Higher utilization. Centralized fleet optimization. The difference between owning a car and accessing mobility becomes meaningful. In that world, Tesla controls both the hardware and the software stack, which allows it to capture far more value per mile than traditional automakers ever could.

Humanoid robotics extends the same logic beyond transportation. The same neural networks used for perception and movement transfer into physical labor. Warehouses. Factories. Logistics. Repetitive tasks that struggle to find human labor become software problems. When labor becomes programmable, productivity scales in a way that traditional economics never modeled well.

Energy and storage quietly reinforce the story. Autonomous systems require stable power. Data centers require storage. Grid instability becomes a tailwind for distributed energy solutions. Tesla sits across all of these layers, which allows internal synergies that competitors cannot easily replicate.

The financial profile reflects a business transitioning phases. Capital intensity remains high today because Tesla is building optionality. Over time, software margins begin to dominate. Once autonomy and robotics mature, revenue decouples from unit sales. That shift matters because software driven cash flows behave very differently than manufacturing margins.

There are real risks. Execution must remain strong. Regulation can slow deployment. Public sentiment will swing violently. The stock will remain volatile because the future arrives unevenly. None of that changes the long arc.

Over a ten year horizon, Tesla benefits from a simple force. Anything that can be automated eventually will be. Tesla owns one of the most advanced real world automation platforms on the planet. Cars happen to be the first expression of that capability, not the final one.

This is why Tesla belongs on this list. It sits at the intersection of autonomy, robotics, energy, and software. Few companies have that many compounding vectors running at the same time.

Nvidia \$NVDA

The Hardware Backbone of AI

Nvidia owns the most important real estate in artificial intelligence. Compute. Every serious AI workload begins and ends with silicon, and Nvidia sits at the center of that demand curve. Models get headlines. Hardware pays the bills.

What makes Nvidia powerful is not a single chip. It is the full platform. GPUs. Networking. Software libraries. Developer tools. CUDA. Enterprise support. Once a customer builds on this stack, moving away becomes expensive, slow, and risky. That creates inertia, and inertia creates pricing power.

AI workloads scale aggressively. Training larger models requires exponentially more compute. Inference at scale requires massive parallel processing. Data centers do not buy one chip. They build entire systems around a vendor they trust. Nvidia benefits from this behavior because its ecosystem reduces friction across the entire deployment cycle.

The company also benefits from neutrality. Cloud providers. Startups. Governments. Enterprises. Everyone buys compute. Nvidia does not need to pick winners at the model level. If demand for AI grows, demand for compute grows. That linkage has proven durable across cycles.

Networking has become a critical amplifier. As clusters grow, the ability to move data efficiently becomes as important as raw compute. Nvidia's expansion into high speed interconnects increases its share of wallet while deepening lock in. The more pieces Nvidia supplies, the harder it becomes to replace any single one.

Margins reflect this position. High gross margins signal scarcity and differentiation. As long as compute remains constrained and demand remains strong, Nvidia captures value upstream where it belongs. Over time, even as competition increases, the installed base and software layer act as stabilizers.

Risks exist. Customers want alternatives. Governments care about supply chains. Competition will not stand still. These forces may compress margins at times. They do not erase the central role Nvidia plays in the AI economy.

Over a 10 year horizon, Nvidia benefits from a simple reality. Intelligence at scale requires massive compute. That requirement does not disappear. It grows. Nvidia built the default platform early, and defaults tend to persist.

This is why Nvidia earns its place here. It supplies the physical substrate on which the AI economy runs. As long as machines keep learning, Nvidia keeps selling shovels.

Microsoft \$MSFT

The Real Estate Layer of AI

Microsoft owns something far more powerful than a single product cycle. It owns positioning. Azure sits underneath governments, banks, defense contractors, hospitals, and global enterprises that value stability over experimentation. In the AI era, that positioning turns cloud infrastructure into digital real estate. Whoever owns the land collects rent as activity increases.

AI workloads amplify this advantage. Training models, running inference, storing sensitive data, and deploying tools across large organizations require compliance, security, uptime, and global reach. Azure already lives inside those requirements. Customers expand usage because the safest path forward usually involves building more deeply into what already works. Migration risk remains high, which strengthens retention and pricing power over time.

Distribution is where Microsoft quietly dominates. AI gets embedded into products that hundreds of millions of people already use every day. Office, Teams, GitHub, Dynamics, Windows. Adoption happens through workflow integration rather than behavior change. That reduces friction and accelerates monetization. Enterprises do not need to rethink how they operate. They enhance existing systems with intelligence.

The enterprise moat deepens through identity, security, and access control. Once a company standardizes on Microsoft across cloud, productivity, and developer tools, the ecosystem becomes cohesive. Each layer reinforces the others. The cost of replacing that stack rises every year, especially as AI touches more sensitive decisions and data flows.

Capital allocation strengthens the thesis. Microsoft generates enormous free cash flow and reinvests it aggressively into data centers, custom silicon, and infrastructure without stressing the balance sheet. Scale becomes an advantage rather than a liability. Smaller competitors struggle to match spending intensity. Larger customers prefer vendors that can keep up with long term demand.

AI does not require Microsoft to win every technical benchmark. It requires Microsoft to remain the default environment where enterprises deploy intelligence safely. That role aligns perfectly with its historical strengths. Reliability. Enterprise trust. Long term contracts. Global compliance.

There are risks. Cloud competition remains intense. Regulatory scrutiny will increase. Capital expenditure cycles can compress margins temporarily. None of these undermine the core position. They shape the pace, not the direction.

Over a ten year horizon, Microsoft benefits from enterprises choosing familiarity, reliability, and scale as AI becomes mission critical. When intelligence moves from experimentation to infrastructure, landlords tend to win.

Microsoft earns its place on this list because it monetizes growth through ownership rather than hype. As AI workloads expand, the rent gets paid every month.

CrowdStrike

The Security Layer of the AI Economy

As systems become more automated, the cost of failure rises dramatically. When decisions are made by software and executed at machine speed, a single breach can ripple across an entire organization in minutes. CrowdStrike sits directly at that pressure point. It protects the endpoints, identities, and cloud workloads that keep modern businesses running.

CrowdStrike's strength comes from architecture. The platform was built cloud native from day one. Every endpoint feeds data into a central intelligence layer. That data trains detection models continuously. Each new customer improves visibility. Each new attack improves defense. Scale strengthens the product rather than diluting it, which is a rare trait in software.

AI accelerates both the number and sophistication of attacks. Automated systems create new surfaces to exploit. Identity becomes the new perimeter as users, machines, and applications interact constantly. CrowdStrike adapts well to this shift because it focuses on behavior rather than static rules. It observes how systems act over time, which allows it to catch threats that never appeared before. This approach becomes more valuable as attackers also adopt automation.

Trust sits at the center of the business. Security teams choose platforms they believe will not fail under pressure. Once that trust is earned, switching feels irresponsible rather than opportunistic. CrowdStrike benefits from this dynamic. Customers often

start with endpoint protection and then expand into identity, cloud, and threat intelligence. Each added module tightens the relationship and raises the cost of leaving.

The financial model reflects that durability. Recurring revenue dominates. Gross margins remain strong. As the platform grows, incremental costs rise far more slowly than revenue. This creates operating leverage that compounds quietly over time, often without dramatic headlines.

There are real risks. Competition remains intense. Expectations remain high. Any visible failure would be punished quickly by customers and markets alike. These are not trivial concerns. They are part of operating in a mission critical category.

Over a ten year horizon, the environment continues to tilt in CrowdStrike's favor. More software. More automation. More interconnected systems. More consequences when something breaks. In that world, security becomes foundational rather than optional.

CrowdStrike earns its place on this list because it protects the nervous system of the digital economy. As AI accelerates everything else, demand for reliable, adaptive security only grows. That creates a long runway built on necessity rather than hype.

Constellation Energy \$CEG

Powering the AI Economy

Every conversation about AI eventually runs into a physical constraint. Electricity. Models do not train themselves on optimism. Data centers do not run on narratives. They run on reliable power, delivered every second of every day, at massive scale. Constellation Energy sits directly at that constraint.

Constellation owns the largest nuclear fleet in the United States. Nuclear power provides constant baseload energy without interruption. That matters because AI workloads cannot tolerate volatility. Training clusters, inference systems, and mission critical cloud services require stability measured in seconds, not averages across a day. When uptime becomes non negotiable, the value of dependable power rises sharply.

As data center demand accelerates, power availability becomes a bottleneck. Building new generation takes years. Upgrading grids takes longer. That scarcity shifts leverage toward existing producers with proven assets already connected to the grid. Constellation benefits from this dynamic without needing speculative expansion. The assets are built. The demand is coming to them.

Nuclear also fits the policy environment. Governments want clean energy without sacrificing reliability. Nuclear satisfies both requirements. That alignment creates regulatory support rather than resistance. Long term contracts add visibility. Customers value certainty when planning multi billion dollar infrastructure projects, and Constellation can offer it.

Pricing power follows necessity. When energy becomes strategic rather than commoditized, suppliers gain leverage. Data center operators care more about reliability than marginal price differences. Losing power costs far more than paying a premium for stability. That reality supports durable cash flows.

Financially, the business reflects infrastructure economics. Predictable revenue. Long duration assets. Cash flows that support reinvestment and shareholder returns. This is not a company chasing explosive growth. It compounds through relevance.

There are risks. Public sentiment around nuclear can shift. Regulatory frameworks can change. Capital requirements remain high. These factors affect pacing, not purpose. The role nuclear plays in an electrified, automated economy remains intact.

Over a ten year horizon, energy demand grows alongside compute demand. AI increases electricity consumption structurally, not cyclically. Power becomes a limiting factor for digital expansion.

Constellation earns its place here because it converts electrons into certainty. As the digital world scales, the businesses that keep the lights on quietly gain leverage. This is foundational infrastructure doing exactly what long term investing rewards.

Vertiv \$VRT

Cooling the Digital World

Compute creates heat. As AI workloads scale, heat becomes one of the most immediate physical limits inside data centers. Chips can be world class and power can be abundant, but without effective cooling the entire system throttles. Vertiv sits exactly at that constraint.

Vertiv provides the infrastructure that keeps data centers operating within safe thermal ranges. Precision cooling. Power management. Monitoring systems. These are not optional accessories. They are mission critical components that determine uptime, efficiency, and hardware lifespan. As racks get denser and chips draw more power, thermal management moves from background concern to central design priority.

AI accelerates this shift. Training clusters generate intense heat loads in concentrated spaces. Traditional cooling approaches struggle to keep up. Vertiv's solutions are built for high density environments where failure costs millions per hour. When customers design new data centers, cooling architecture gets planned early, not added later. That embeds Vertiv deeply into the build cycle.

The business benefits from long replacement cycles and recurring service revenue. Once installed, infrastructure remains in place for years. Maintenance, upgrades, and monitoring create ongoing relationships rather than one time transactions. Switching suppliers mid lifecycle introduces risk most operators prefer to avoid. That behavior supports retention and pricing discipline.

Energy efficiency adds another tailwind. Cooling consumes a large share of data center power. Better thermal management reduces operating costs and improves sustainability metrics. As electricity prices rise and energy scrutiny increases, efficiency becomes a financial lever rather than a marketing line.

Vertiv operates quietly. It does not sell dreams. It sells reliability. That profile often leads to underappreciation during hype cycles. Over long periods, it leads to steady compounding as infrastructure spending grows.

There are risks. Capital spending can be cyclical. Competition exists. Execution matters. These factors affect timing rather than relevance. As long as compute density increases, cooling remains essential.

Over a ten year horizon, AI pushes physical limits harder each year. Data centers evolve from buildings into highly engineered machines. Vertiv supplies the systems that keep those machines alive.

Vertiv earns its place on this list because it monetizes necessity. When digital ambition meets physical reality, companies that solve the bottlenecks tend to compound quietly and persistently.

Google \$GOOG

Cloud, Models, and Custom Silicon

Google occupies a unique position in the AI stack because it controls multiple critical layers at once. Cloud infrastructure. Proprietary models. Custom silicon. That combination gives Google flexibility and leverage that few companies can match over a ten year horizon.

Google Cloud has matured into a serious enterprise platform. It supports large scale data workloads, advanced analytics, and AI deployment for companies that care about performance and cost efficiency. As AI moves deeper into production environments, customers look for platforms that can handle massive data volumes with reliability. Google's internal experience running search, video, and advertising at global scale translates directly into that capability.

Models matter, but deployment matters more. Gemini strengthens Google's ability to offer integrated AI services across search, cloud, and productivity tools. The value here does not come from winning every benchmark. It comes from embedding intelligence into products used by billions of people and thousands of enterprises. Distribution remains one of Google's strongest advantages.

The most underappreciated asset sits beneath the surface. TPUs. Google built custom silicon to optimize AI workloads long before it became fashionable. These chips reduce cost, improve efficiency, and give Google control over its own compute destiny. TPUs work alongside GPUs rather than replacing them. That flexibility allows Google to balance performance, cost, and scale depending on the workload.

Owning custom silicon also changes economics. When compute becomes cheaper internally, Google can reinvest savings into pricing, capacity, or product improvement. That creates a competitive loop where efficiency fuels growth. Over time, this matters more than short term performance comparisons.

Advertising cash flows support everything else. Search and video generate enormous free cash flow that funds cloud expansion, model development, and infrastructure without stressing the balance sheet. This internal subsidy allows Google to invest patiently while others depend on external capital.

There are risks. Regulation remains a constant overhang. Cloud competition stays intense. Product decisions do not always land cleanly. These factors create volatility. They do not erase Google's structural advantages.

Over a ten year window, Google benefits from controlling both demand and supply of intelligence. It builds models, runs them at scale, and distributes them through products people already use daily.

Google earns its place on this list because it compounds through integration. As AI becomes embedded everywhere, companies that own multiple layers of the stack gain resilience and optionality. That combination gives Google long runway and strategic depth.

Broadcom \$AVGO

The Toll Roads of Enterprise Technology

Broadcom operates in parts of the technology stack that most investors rarely talk about, yet almost every large enterprise depends on. It supplies the connective tissue that allows modern systems to function at scale. Networking chips. Custom silicon. Infrastructure software. These are the toll roads that data, compute, and workloads must pass through every day.

The core strength of Broadcom comes from positioning rather than hype. Its products are embedded deep inside data centers, telecom networks, and enterprise systems. Once installed, they tend to stay in place for long periods. Replacement introduces risk.

Risk introduces hesitation. That hesitation turns into pricing power and long duration revenue.

AI increases demand for everything Broadcom touches. More compute means more networking. Larger clusters require faster interconnects. Data must move efficiently between chips, racks, and buildings. Broadcom benefits because it supplies components that scale with system complexity rather than end user excitement. As architectures grow more demanding, reliability matters more than novelty.

Custom silicon adds another layer of durability. Hyperscalers want differentiation and cost control. Broadcom designs chips tailored to specific workloads, locking in long term relationships. These partnerships tend to last for many years because redesign cycles are slow and expensive. That creates visibility rarely seen in technology businesses.

The software side reinforces the model. Enterprise infrastructure software generates recurring revenue with high margins and low churn. Customers rely on these tools to manage, secure, and optimize critical systems. Once software becomes part of daily operations, switching feels disruptive. That behavior supports steady cash generation.

Broadcom's approach to capital allocation strengthens the story. Management focuses on cash flow, disciplined investment, and shareholder returns. Growth does not need to be explosive to be effective. Durability combined with scale produces compounding that looks boring quarter to quarter and powerful decade to decade.

There are risks. Customer concentration exists. Cycles affect capital spending. Integration of acquisitions must be handled carefully. These factors influence volatility. They do not undermine the structural role Broadcom plays.

Over a ten year horizon, enterprise technology continues to expand quietly behind the scenes. Data volumes grow. Connectivity demands rise. Customization becomes more valuable. Broadcom sits in the middle of these trends, collecting rent as traffic increases.

Broadcom earns its place on this list because it monetizes necessity rather than excitement. When digital systems grow more complex, the owners of the toll roads tend to compound steadily and persistently.

The Consumer Ecosystem Compounder

Apple occupies a position no other technology company has been able to replicate. It owns the consumer interface layer at global scale. Over 2 billion active devices live inside its ecosystem. That installed base creates a locked in audience that turns incremental improvements into massive cash flow over time.

The real strength of Apple sits in services and ecosystem control rather than headline innovation. Hardware sales establish the relationship. Software, subscriptions, payments, storage, media, and accessories monetize it year after year. Once users are inside the ecosystem, friction keeps them there. Switching devices often means switching habits, data, and daily routines. Most people simply do not bother.

This dynamic produces one of the most powerful business models ever built. High margin services grow on top of a hardware foundation that refreshes predictably. Revenue becomes recurring. Cash flows become stable. That stability allows Apple to allocate capital aggressively without threatening the core business.

Buybacks play a critical role. Apple returns enormous amounts of capital to shareholders while continuing to invest in product development and supply chain resilience. Reducing share count quietly amplifies per share growth even during periods when headline revenue growth slows. This is compounding in its most disciplined form.

AI does not need to redefine Apple's identity to matter. Intelligence enhances the ecosystem by improving usability, personalization, and efficiency. Features layered into devices and operating systems strengthen user attachment. Apple benefits even when it moves deliberately because its advantage comes from distribution rather than speed.

Privacy and trust reinforce the moat. Consumers increasingly care about how data gets used. Apple positions itself as a steward rather than an extractor. That perception supports brand loyalty and pricing power, especially in a world where digital trust becomes scarcer.

There are risks. Hardware cycles can mature. Regulatory scrutiny continues. Growth rates will never look explosive. These realities are already reflected in expectations.

What matters is durability. Few companies have shown the ability to monetize an installed base this large with such consistency.

Over a ten year horizon, Apple benefits from owning the front door to digital life. As services expand and devices remain central to daily activity, cash flows continue to compound quietly.

Apple earns its place on this list because it turns consumer habit into shareholder return. No reinvention required. Just discipline, scale, and time doing the work.

Amazon \$AMZN

The Infrastructure Spine of the Digital Economy

Amazon built one of the most misunderstood compounding machines in modern business. Many still view it through the lens of retail margins and logistics headaches. The real value sits underneath that surface in infrastructure, scale, and optionality that continues to expand over time.

AWS is the centerpiece. It functions as critical infrastructure for the internet economy. Startups, enterprises, governments, and global platforms rely on it to store data, run applications, and deploy intelligence. As AI workloads grow, cloud demand grows with them. Training, inference, storage, and networking all scale together. AWS captures value at every step of that expansion.

The power of AWS comes from trust and breadth. Customers do not choose a cloud provider lightly. Migration is costly. Downtime is unacceptable. Once workloads are deployed, usage tends to expand organically. New services get layered on top of existing ones. That behavior produces recurring revenue and long duration relationships.

Retail strengthens the ecosystem rather than weakening it. Logistics, fulfillment, and last mile delivery create a physical network that is extremely difficult to replicate. That network supports Prime, advertising, and third party sellers. Advertising has emerged as a high margin business riding on top of commerce traffic, adding another compounding cash flow stream.

Operational discipline has improved meaningfully. After years of heavy investment, efficiency gains flow through the income statement. When growth reaccelerates, margins expand faster than revenue. This operating leverage often surprises investors who focus too narrowly on retail economics.

Amazon also benefits from optionality. Logistics can support healthcare. Data can support advertising. Cloud infrastructure can support AI services. The company continuously converts internal capability into external revenue. That ability comes from scale and patience rather than short term optimization.

There are risks. Regulation remains a concern. Competition in cloud remains intense. Capital intensity never disappears. These factors affect sentiment and valuation at times. They do not erase the underlying compounding engine.

Over a ten year horizon, digital activity continues to migrate to the cloud. Commerce continues to move online. Advertising follows attention. Infrastructure owners capture value at each layer.

Amazon earns its place on this list because it monetizes growth in the digital economy from multiple directions at once. When systems scale, platforms that sit underneath them tend to compound relentlessly. Amazon remains one of the strongest examples of that dynamic in action.

Oracle \$ORCL

Locking In Enterprise AI Workloads

Oracle has spent years being underestimated, largely because it does not chase attention. That has turned out to be an advantage. While others focused on consumer narratives, Oracle focused on becoming indispensable to enterprises that run the world's most critical systems. In the AI era, that positioning becomes more valuable, not less.

Oracle owns deep relationships with enterprises that cannot afford disruption. Databases. Financial systems. Supply chains. Telecom billing. Government records. These workloads are sticky by nature. They involve sensitive data, regulatory

oversight, and operational risk. Moving them is expensive and often avoided entirely. Oracle understands this psychology and builds around it.

OCI strengthens that grip. Oracle Cloud Infrastructure is designed specifically for high performance, low latency, and cost efficient enterprise workloads. AI training and inference fit naturally into this environment when data already lives inside Oracle systems. Instead of exporting data to another cloud, enterprises increasingly prefer to bring compute to the data. That preference benefits Oracle directly.

AI adoption inside enterprises does not look like experimentation. It looks like integration. Forecasting. Fraud detection. Resource optimization. Decision support. Oracle's strength lies in embedding intelligence into existing workflows rather than forcing companies to reinvent how they operate. That approach accelerates adoption because it reduces organizational resistance.

Pricing discipline matters here. Oracle competes aggressively on cost while maintaining reliability. For CIOs under pressure to control budgets, predictable pricing combined with performance is attractive. Once workloads are deployed, expansion tends to follow usage patterns rather than marketing cycles.

The business model reflects long term thinking. High recurring revenue. Long contracts. Strong visibility. Cash flows support investment and shareholder returns without chasing aggressive growth optics. This creates compounding that looks slow on the surface and powerful over time.

There are risks. Oracle faces strong competition from hyperscalers. Brand perception still lags technical reality. Execution must remain consistent. These factors influence pace rather than direction.

Over a ten year horizon, enterprises move AI deeper into core operations. They prioritize vendors that understand mission critical systems and regulatory environments. Oracle fits that requirement naturally.

Oracle earns its place on this list because it locks in the workloads that matter most. As AI becomes embedded into enterprise decision making, the platforms closest to the data quietly gain leverage. Oracle is built for exactly that outcome.

ARM: \$ARM

The Instruction Set Foundation

ARM sits at one of the deepest layers of modern computing. It does not sell finished products. It licenses the instruction set architecture that tells chips how to operate. That position gives ARM influence over everything built on top of it, from smartphones to servers to edge devices to emerging AI hardware.

The power of ARM comes from ubiquity. Most of the world's mobile devices run on ARM based designs. That dominance was earned through efficiency. Low power consumption. High performance per watt. As computing spreads into more places and more devices, efficiency matters more than raw speed. This trend works directly in ARM's favor.

AI extends the opportunity. Models do not live only in massive data centers. They move to phones, cars, sensors, factories, and edge systems. ARM architecture fits these environments naturally. It allows designers to build specialized chips optimized for specific workloads while remaining compatible with a massive ecosystem of software tools.

The licensing model creates leverage. ARM collects royalties on volume rather than carrying manufacturing risk. As more devices ship, revenue grows without proportional cost increases. This produces operating leverage that compounds over time. Customers invest heavily in ARM based designs, which deepens commitment and reduces the likelihood of switching architectures.

ARM also benefits from neutrality. It does not compete directly with its customers. Chip designers, cloud providers, and device manufacturers all build on top of ARM without fearing platform conflict. That trust supports adoption across industries and geographies.

Data centers represent a growing frontier. Power efficiency has become a central concern as AI workloads expand. ARM based server chips offer an alternative that prioritizes performance per watt. As operators balance energy costs and compute needs, architectural diversity becomes attractive. ARM stands to benefit from this shift without needing to displace every existing solution.

There are risks. Customer concentration matters. Competition from other architectures remains. Execution in new markets must continue to prove itself. These factors affect adoption speed rather than long term relevance.

Over a ten year horizon, computing continues to spread everywhere. Devices become smarter. Systems become more distributed. Power efficiency becomes a constraint. Instruction sets that scale across environments gain value.

ARM earns its place on this list because it owns a foundational layer that everything else builds upon. When an architecture becomes a standard, time works in its favor. ARM has already crossed that threshold.

AMD: \$AMD

Alternative Compute at Scale

AMD plays a critical role in the AI and data center ecosystem by providing credible alternatives where concentration risk becomes unacceptable. Large customers rarely want dependency on a single supplier, especially when compute sits at the center of strategic infrastructure. AMD benefits directly from that reality.

The company has rebuilt its reputation around execution. Over the past several years, AMD delivered consistent architectural improvements across CPUs and accelerators. Performance gains came alongside efficiency gains. That balance matters in data centers where power, cooling, and total cost of ownership drive decisions as much as raw speed.

In AI, AMD offers customers choice. Training and inference workloads vary. Some prioritize peak performance. Others prioritize cost efficiency or power usage. AMD's accelerators allow hyperscalers and enterprises to diversify deployments without rewriting entire software stacks. That flexibility lowers adoption friction and increases willingness to experiment at scale.

CPUs remain a major strength. As AI workloads expand, supporting infrastructure matters more. Data preprocessing. Orchestration. Memory management. AMD's server CPUs continue to gain share because they offer strong performance per dollar. These wins compound quietly as customers standardize platforms across fleets.

The strategic value of AMD increases as markets grow. In trillion dollar compute markets, being the second major supplier still creates enormous opportunity. Customers reward vendors that show reliability, roadmap clarity, and pricing discipline. AMD fits that profile.

The business model reflects this positioning. Margins expand as mix improves. Operating leverage builds as volumes rise. The company does not need to dominate every segment to compound effectively. It needs to remain relevant, competitive, and trusted.

There are risks. Competition remains intense. Nvidia sets a high bar in accelerators. Execution must remain consistent. These realities create volatility. They do not erase the role AMD plays in balancing the ecosystem.

Over a ten year horizon, demand for compute continues to grow across AI, cloud, and enterprise workloads. Concentration risk becomes a strategic concern. Alternative suppliers gain leverage simply by existing at scale.

AMD earns its place on this list because it supplies optionality in a world that increasingly values resilience. When compute demand explodes, customers rarely choose dependency. They choose flexibility. AMD benefits from that instinct.

ASML \$ASML

The Ultimate Bottleneck

ASML occupies one of the rarest positions in global industry. It controls a capability that cannot be substituted, bypassed, or accelerated through competition. Advanced semiconductor manufacturing depends on extreme ultraviolet lithography, and ASML is the only company on earth that can build and deliver these systems.

Every leading edge chip begins here. Logic chips. Memory. AI accelerators. CPUs. GPUs. If a manufacturer wants to move to smaller nodes with higher density and better performance, ASML sits directly in the path. This creates a structural bottleneck that turns technological leadership into long duration pricing power.

EUV systems are among the most complex machines ever built. They combine optics, lasers, software, materials science, and precision engineering at a level that took decades to achieve. Replicating this capability would require enormous capital, time, talent, and coordination across entire supply chains. That barrier protects ASML in a way few technology companies ever experience.

AI strengthens this position. As demand for compute rises, chipmakers race to improve efficiency and performance. Smaller nodes allow more transistors, lower power usage, and better economics at scale. That race feeds directly into ASML's order book. More AI means more chips. More advanced chips mean more reliance on EUV.

The business model reflects scarcity. Each system sells for enormous sums. Service and maintenance create recurring revenue over the lifespan of the machines. Customers plan years in advance because production timelines depend on ASML deliveries. That visibility supports disciplined growth rather than reactive expansion.

Geopolitics adds another layer. Governments recognize semiconductor manufacturing as strategic infrastructure. Export controls, supply chain security, and national interests all converge on lithography. ASML operates at the center of this tension while remaining indispensable to allies. That position reinforces its relevance even as the world fragments.

There are risks. Cyclicalities in semiconductor spending affects timing. Political pressure creates complexity. Customer concentration remains high. These factors influence volatility. They do not weaken the bottleneck.

Over a ten year horizon, the direction remains clear. Demand for compute rises. Chips get more advanced. Manufacturing complexity increases. Bottlenecks gain value.

ASML earns its place on this list because it controls the narrowest choke point in modern technology. When an entire industry must pass through a single gate, the owner of that gate tends to compound quietly and relentlessly.

Conclusion: The Full AI Value Chain

Looking at these companies individually only tells part of the story. The real power emerges when you see how they connect. AI is not a single product or a single winner. It is a system made of interlocking layers, each dependent on the others. These businesses sit at those layers.

It starts with energy. Without reliable electricity, nothing else exists. Constellation supplies the stable power that data centers require to run continuously. Power feeds compute. Compute generates heat. Vertiv keeps those systems operating within physical limits. Cooling becomes as critical as the chips themselves once density rises.

On top of that physical foundation sits silicon. ASML enables the most advanced chips to be manufactured. Nvidia and AMD supply the compute that trains and runs models. ARM provides the instruction set that allows intelligence to move efficiently across devices, from data centers to the edge. Broadcom connects everything through networking and custom silicon that moves data where it needs to go.

Cloud platforms bring these components together. Microsoft, Amazon, Google, and Oracle provide the environments where AI workloads live. They offer scale, security, compliance, and global reach. Enterprises deploy intelligence inside these platforms because moving workloads repeatedly is expensive and risky. Once deployed, usage tends to expand rather than retreat.

Software turns capability into outcomes. Palantir sits above models and infrastructure, translating data into decisions. CrowdStrike protects the entire system from intrusion and failure. Apple owns the consumer interface layer where intelligence meets daily life. Tesla extends intelligence into the physical world through autonomy and robotics.

Each layer reinforces the others. More AI increases demand for power. More power enables more compute. More compute requires better cooling and networking. More intelligence raises the value of security and orchestration. This feedback loop compounds over time.

What matters most is that none of these businesses rely on a single narrow outcome. They benefit from growth across the system. When one layer accelerates, the others follow. This reduces dependency on any single prediction about how AI evolves.

Over a ten year horizon, systems thinking beats stock picking. Investors who understand how value flows through the stack gain patience during volatility and clarity during hype cycles. This list was built around that principle.

Own the layers.

Own the bottlenecks. Own the platforms that others depend on.

When the system grows, the owners of the system tend to win.

© 2025 Tom. All rights reserved.

This document is for informational purposes only and does not constitute financial advice.